

# Assignment 4.1 Mood Detection with OpenCV

April 13, 2025

Technological Institute of the Philippines	Quezon City - Computer Engineering
Course Code:	CPE 313
Course Title:	Advanced Machine Learning and Deep Learning
2nd Semester	AY 2024-2025
<b>ACTIVITY NO.</b>	Assignment 4.1 Mood Detection with OpenCV
<b>Name</b>	Base, Angelo P.
<b>Section</b>	CPE32S1
<b>Date Performed:</b>	06/04/2025
<b>Date Submitted:</b>	13/04/2025
<b>Instructor:</b>	Dr. Alonica Villanueva

```
[2]: import os
import cv2
import numpy as np

import pickle
from tqdm import tqdm

from tensorflow.keras.preprocessing.image import ImageDataGenerator, img_to_array
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Sequential, load_model
from tensorflow.keras.layers import Flatten, Dense, Dropout
from tensorflow.keras.optimizers import Adam

from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix

import itertools

from deepface import DeepFace
import matplotlib.pyplot as plt
# from google.colab.patches import cv2_imshow
```

### 0.0.1 Face Recognition

```
[ ]: dataset_path = "custom_facial_expression_dataset/train"
my_face_folders = ["angry", "happy", "sad", "confused"]
others_folder = "custom_facial_expression_dataset/others"

embeddings = []
labels = []

# Process my_face_folders images
for folder in my_face_folders:
    folder_path = os.path.join(dataset_path, folder)
    for img_name in os.listdir(folder_path):
        img_path = os.path.join(folder_path, img_name)
        try:
            embedding = DeepFace.represent(img_path,
↪model_name="Facenet")[0]["embedding"]
            embeddings.append(embedding)
            labels.append(1)
        except:
            print(f"Skipping image: {img_path}")

# Process others images
for img_name in os.listdir(others_folder):
    img_path = os.path.join(others_folder, img_name)
    try:
        embedding = DeepFace.represent(img_path,
↪model_name="Facenet")[0]["embedding"]
        embeddings.append(embedding)
        labels.append(0)
    except:
        print(f"Skipping image: {img_path}")

embeddings = np.array(embeddings)
labels = np.array(labels)

with open("pkl/face_embeddings.pkl", "wb") as f:
    pickle.dump((embeddings, labels), f)
print("Saved!")
```

```
Skipping image: custom_facial_expression_dataset/train\angry\Training_10.jpg
Skipping image: custom_facial_expression_dataset/train\angry\Training_110.jpg
Skipping image: custom_facial_expression_dataset/train\angry\Training_116.jpg
Skipping image: custom_facial_expression_dataset/train\angry\Training_117.jpg
Skipping image: custom_facial_expression_dataset/train\angry\Training_119.jpg
Skipping image: custom_facial_expression_dataset/train\angry\Training_12.jpg
Skipping image: custom_facial_expression_dataset/train\angry\Training_120.jpg
Skipping image: custom_facial_expression_dataset/train\angry\Training_121.jpg
```







[illegible]

Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_46.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_47.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_49.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_5.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_50.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_53.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_54.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_55.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_56.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_58.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_59.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_6.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_64.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_67.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_68.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_69.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_8.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_81.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_82.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_90.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_91.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_92.jpg  
 Skipping image: custom\_facial\_expression\_dataset/train/confused/Training\_93.jpg  
 Skipping image: custom\_facial\_expression\_dataset/others/20230604\_010504.jpg  
 Saved!

```

[ ]: X_train, X_test, y_train, y_test = train_test_split(embeddings, labels,
    ↪test_size=0.2, random_state=42)

svm_model = SVC(kernel='linear', probability=True)
svm_model.fit(X_train, y_train)

y_pred = svm_model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Face Recognition Accuracy: {accuracy * 100:.2f}%")

with open("pkl/face_recognition_model.pkl", "wb") as f:
    pickle.dump(svm_model, f)
print("Saved!")
  
```

Face Recognition Accuracy: 98.14%  
 Saved!

## 0.1 Mood Detection

### 0.1.1 Load Dataset

```
[8]: TRAIN_DIR = ('custom_facial_expression_dataset/train/')
TEST_DIR = ('custom_facial_expression_dataset/test/')
VAL_DIR = ('custom_facial_expression_dataset/validation/')
```

```
[9]: def load_data(dir_path, IMG_SIZE):

    X = []
    y = []
    i = 0
    labels = dict()
    for path in tqdm(sorted(os.listdir(dir_path))):
        if not path.startswith('.'):
            labels[i] = path
            for file in os.listdir(dir_path + path):
                if not file.startswith('.'):
                    img = cv2.imread(dir_path + path + '/' + file)
                    img = img.astype('float32') / 255
                    resized = cv2.resize(img, IMG_SIZE, interpolation = cv2.
↪INTER_AREA)

                    X.append(resized)
                    y.append(i)

            i += 1
    X = np.array(X)
    y = np.array(y)
    print(f'{len(X)} images loaded from {dir_path} directory.')
    return X, y, labels

IMG_SIZE = 48
```

```
[10]: X_train, y_train, train_labels = load_data(TRAIN_DIR, (IMG_SIZE, IMG_SIZE))
X_test, y_test, test_labels = load_data(TEST_DIR, (IMG_SIZE, IMG_SIZE))
X_val, y_val, val_labels = load_data(VAL_DIR, (IMG_SIZE, IMG_SIZE))
```

```
0%|          | 0/4 [00:00<?, ?it/s]100%|          | 4/4 [00:26<00:00,
6.72s/it]

800 images loaded from custom_facial_expression_dataset/train/ directory.
100%|          | 4/4 [00:03<00:00, 1.01it/s]

120 images loaded from custom_facial_expression_dataset/test/ directory.
100%|          | 4/4 [00:04<00:00, 1.08s/it]

120 images loaded from custom_facial_expression_dataset/validation/ directory.
```



```
[11]: train_labels
```

```
[11]: {0: 'angry', 1: 'confused', 2: 'happy', 3: 'sad'}
```

```
[12]: test_labels
```

```
[12]: {0: 'angry', 1: 'confused', 2: 'happy', 3: 'sad'}
```

```
[13]: val_labels
```

```
[13]: {0: 'angry', 1: 'confused', 2: 'happy', 3: 'sad'}
```

```
[14]: from keras.utils.np_utils import to_categorical

y_train = to_categorical(y_train, num_classes=4)
y_train.shape
```

```
[14]: (800, 4)
```

```
[15]: y_test = to_categorical(y_test, num_classes=4)
y_test.shape
```

```
[15]: (120, 4)
```

```
[16]: y_val = to_categorical(y_val, num_classes=4)
y_val.shape
```

```
[16]: (120, 4)
```

### 0.1.2 Creating Model

```
[122]: conv_base = VGG16(weights='imagenet',
                        include_top=False,
                        input_shape=(IMG_SIZE, IMG_SIZE, 3)
)

for layer in conv_base.layers[:15]:
    layer.trainable = False

for layer in conv_base.layers[15:]:
    layer.trainable = True

conv_base.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
input_13 (InputLayer)	[(None, 48, 48, 3)]	0

block1_conv1 (Conv2D)	(None, 48, 48, 64)	1792
block1_conv2 (Conv2D)	(None, 48, 48, 64)	36928
block1_pool (MaxPooling2D)	(None, 24, 24, 64)	0
block2_conv1 (Conv2D)	(None, 24, 24, 128)	73856
block2_conv2 (Conv2D)	(None, 24, 24, 128)	147584
block2_pool (MaxPooling2D)	(None, 12, 12, 128)	0
block3_conv1 (Conv2D)	(None, 12, 12, 256)	295168
block3_conv2 (Conv2D)	(None, 12, 12, 256)	590080
block3_conv3 (Conv2D)	(None, 12, 12, 256)	590080
block3_pool (MaxPooling2D)	(None, 6, 6, 256)	0
block4_conv1 (Conv2D)	(None, 6, 6, 512)	1180160
block4_conv2 (Conv2D)	(None, 6, 6, 512)	2359808
block4_conv3 (Conv2D)	(None, 6, 6, 512)	2359808
block4_pool (MaxPooling2D)	(None, 3, 3, 512)	0
block5_conv1 (Conv2D)	(None, 3, 3, 512)	2359808
block5_conv2 (Conv2D)	(None, 3, 3, 512)	2359808
block5_conv3 (Conv2D)	(None, 3, 3, 512)	2359808
block5_pool (MaxPooling2D)	(None, 1, 1, 512)	0

```
=====
Total params: 14,714,688
Trainable params: 7,079,424
Non-trainable params: 7,635,264
-----
```

```
[123]: model = Sequential([
        conv_base,
        Flatten(),
        Dense(512, activation='relu'),
```

```

        Dropout(0.4),
        Dense(4, activation='softmax')
    ])

model.compile(
    optimizer=Adam(learning_rate=1e-5),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

model.summary()

```

Model: "sequential\_13"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 1, 1, 512)	14714688
flatten_13 (Flatten)	(None, 512)	0
dense_26 (Dense)	(None, 512)	262656
dropout_13 (Dropout)	(None, 512)	0
dense_27 (Dense)	(None, 4)	2052
Total params: 14,979,396		
Trainable params: 7,344,132		
Non-trainable params: 7,635,264		

### 0.1.3 Data Augmentation

```

[124]: datagen = ImageDataGenerator(
        rotation_range=20,
        width_shift_range=0.2,
        height_shift_range=0.2,
        shear_range=0.2,
        zoom_range=0.2,
        horizontal_flip=True,
        fill_mode='nearest'
    )

datagen.fit(X_train)

```

### 0.1.4 Train Model

```
[125]: EPOCHS = 50
BATCH_SIZE = 64

history = model.fit(
    datagen.flow(X_train, y_train, batch_size=BATCH_SIZE),
    validation_data = (X_val, y_val),
    epochs=EPOCHS,
    batch_size=BATCH_SIZE,
    verbose=1
)
```

Epoch 1/50

13/13 [=====] - 2s 126ms/step - loss: 1.4976 - accuracy: 0.2587 - val\_loss: 1.3841 - val\_accuracy: 0.2917

Epoch 2/50

13/13 [=====] - 1s 51ms/step - loss: 1.4240 - accuracy: 0.2912 - val\_loss: 1.3688 - val\_accuracy: 0.3250

Epoch 3/50

13/13 [=====] - 1s 52ms/step - loss: 1.3925 - accuracy: 0.3075 - val\_loss: 1.3546 - val\_accuracy: 0.3750

Epoch 4/50

13/13 [=====] - 1s 52ms/step - loss: 1.3382 - accuracy: 0.3600 - val\_loss: 1.3328 - val\_accuracy: 0.4500

Epoch 5/50

13/13 [=====] - 1s 50ms/step - loss: 1.3354 - accuracy: 0.3562 - val\_loss: 1.3178 - val\_accuracy: 0.4417

Epoch 6/50

13/13 [=====] - 1s 52ms/step - loss: 1.2568 - accuracy: 0.4175 - val\_loss: 1.2867 - val\_accuracy: 0.4333

Epoch 7/50

13/13 [=====] - 1s 54ms/step - loss: 1.2503 - accuracy: 0.4375 - val\_loss: 1.2783 - val\_accuracy: 0.4250

Epoch 8/50

13/13 [=====] - 1s 57ms/step - loss: 1.2155 - accuracy: 0.4650 - val\_loss: 1.2501 - val\_accuracy: 0.5167

Epoch 9/50

13/13 [=====] - 1s 54ms/step - loss: 1.1864 - accuracy: 0.4762 - val\_loss: 1.2138 - val\_accuracy: 0.4917

Epoch 10/50

13/13 [=====] - 1s 62ms/step - loss: 1.1407 - accuracy: 0.5125 - val\_loss: 1.2030 - val\_accuracy: 0.4583

Epoch 11/50

13/13 [=====] - 1s 54ms/step - loss: 1.1337 - accuracy: 0.5125 - val\_loss: 1.1679 - val\_accuracy: 0.5417

Epoch 12/50

13/13 [=====] - 1s 52ms/step - loss: 1.0737 - accuracy:

0.5387 - val\_loss: 1.1639 - val\_accuracy: 0.4667  
 Epoch 13/50  
 13/13 [=====] - 1s 64ms/step - loss: 1.0606 - accuracy:  
 0.5425 - val\_loss: 1.1435 - val\_accuracy: 0.4917  
 Epoch 14/50  
 13/13 [=====] - 1s 57ms/step - loss: 1.0351 - accuracy:  
 0.5500 - val\_loss: 1.1662 - val\_accuracy: 0.4583  
 Epoch 15/50  
 13/13 [=====] - 1s 56ms/step - loss: 0.9881 - accuracy:  
 0.5888 - val\_loss: 1.1383 - val\_accuracy: 0.4417  
 Epoch 16/50  
 13/13 [=====] - 1s 52ms/step - loss: 0.9582 - accuracy:  
 0.6175 - val\_loss: 1.0923 - val\_accuracy: 0.5083  
 Epoch 17/50  
 13/13 [=====] - 1s 53ms/step - loss: 0.9173 - accuracy:  
 0.6313 - val\_loss: 1.0813 - val\_accuracy: 0.5000  
 Epoch 18/50  
 13/13 [=====] - 1s 53ms/step - loss: 0.9329 - accuracy:  
 0.6212 - val\_loss: 1.0313 - val\_accuracy: 0.6083  
 Epoch 19/50  
 13/13 [=====] - 1s 54ms/step - loss: 0.9042 - accuracy:  
 0.6250 - val\_loss: 1.0471 - val\_accuracy: 0.5750  
 Epoch 20/50  
 13/13 [=====] - 1s 52ms/step - loss: 0.8628 - accuracy:  
 0.6513 - val\_loss: 1.0306 - val\_accuracy: 0.6000  
 Epoch 21/50  
 13/13 [=====] - 1s 53ms/step - loss: 0.8555 - accuracy:  
 0.6612 - val\_loss: 1.0903 - val\_accuracy: 0.5000  
 Epoch 22/50  
 13/13 [=====] - 1s 54ms/step - loss: 0.8373 - accuracy:  
 0.6525 - val\_loss: 1.0456 - val\_accuracy: 0.5667  
 Epoch 23/50  
 13/13 [=====] - 1s 53ms/step - loss: 0.8597 - accuracy:  
 0.6587 - val\_loss: 1.0393 - val\_accuracy: 0.5750  
 Epoch 24/50  
 13/13 [=====] - 1s 53ms/step - loss: 0.8386 - accuracy:  
 0.6575 - val\_loss: 0.9881 - val\_accuracy: 0.5750  
 Epoch 25/50  
 13/13 [=====] - 1s 53ms/step - loss: 0.7686 - accuracy:  
 0.7013 - val\_loss: 1.0058 - val\_accuracy: 0.5750  
 Epoch 26/50  
 13/13 [=====] - 1s 54ms/step - loss: 0.7505 - accuracy:  
 0.6938 - val\_loss: 1.0124 - val\_accuracy: 0.6000  
 Epoch 27/50  
 13/13 [=====] - 1s 57ms/step - loss: 0.7819 - accuracy:  
 0.6825 - val\_loss: 1.0670 - val\_accuracy: 0.5333  
 Epoch 28/50  
 13/13 [=====] - 1s 53ms/step - loss: 0.7665 - accuracy:

0.6837 - val\_loss: 1.0938 - val\_accuracy: 0.4917  
Epoch 29/50  
13/13 [=====] - 1s 54ms/step - loss: 0.7408 - accuracy:  
0.6975 - val\_loss: 0.9218 - val\_accuracy: 0.6500  
Epoch 30/50  
13/13 [=====] - 1s 52ms/step - loss: 0.7352 - accuracy:  
0.7175 - val\_loss: 0.9751 - val\_accuracy: 0.6083  
Epoch 31/50  
13/13 [=====] - 1s 51ms/step - loss: 0.7218 - accuracy:  
0.7100 - val\_loss: 0.9992 - val\_accuracy: 0.6083  
Epoch 32/50  
13/13 [=====] - 1s 53ms/step - loss: 0.6734 - accuracy:  
0.7262 - val\_loss: 0.9584 - val\_accuracy: 0.6167  
Epoch 33/50  
13/13 [=====] - 1s 54ms/step - loss: 0.6830 - accuracy:  
0.7287 - val\_loss: 1.0450 - val\_accuracy: 0.5667  
Epoch 34/50  
13/13 [=====] - 1s 54ms/step - loss: 0.6708 - accuracy:  
0.7287 - val\_loss: 0.9440 - val\_accuracy: 0.5917  
Epoch 35/50  
13/13 [=====] - 1s 51ms/step - loss: 0.6259 - accuracy:  
0.7738 - val\_loss: 0.9291 - val\_accuracy: 0.6167  
Epoch 36/50  
13/13 [=====] - 1s 54ms/step - loss: 0.6225 - accuracy:  
0.7412 - val\_loss: 1.0581 - val\_accuracy: 0.6333  
Epoch 37/50  
13/13 [=====] - 1s 54ms/step - loss: 0.6341 - accuracy:  
0.7437 - val\_loss: 0.9643 - val\_accuracy: 0.6500  
Epoch 38/50  
13/13 [=====] - 1s 52ms/step - loss: 0.6018 - accuracy:  
0.7725 - val\_loss: 0.9027 - val\_accuracy: 0.6750  
Epoch 39/50  
13/13 [=====] - 1s 52ms/step - loss: 0.5724 - accuracy:  
0.7663 - val\_loss: 0.9509 - val\_accuracy: 0.6583  
Epoch 40/50  
13/13 [=====] - 1s 52ms/step - loss: 0.5979 - accuracy:  
0.7763 - val\_loss: 1.0357 - val\_accuracy: 0.6167  
Epoch 41/50  
13/13 [=====] - 1s 54ms/step - loss: 0.5895 - accuracy:  
0.7775 - val\_loss: 0.9343 - val\_accuracy: 0.6667  
Epoch 42/50  
13/13 [=====] - 1s 54ms/step - loss: 0.5793 - accuracy:  
0.7788 - val\_loss: 0.9821 - val\_accuracy: 0.5833  
Epoch 43/50  
13/13 [=====] - 1s 53ms/step - loss: 0.5775 - accuracy:  
0.7763 - val\_loss: 0.9667 - val\_accuracy: 0.6250  
Epoch 44/50  
13/13 [=====] - 1s 51ms/step - loss: 0.5998 - accuracy:

```

0.7675 - val_loss: 0.9110 - val_accuracy: 0.7083
Epoch 45/50
13/13 [=====] - 1s 52ms/step - loss: 0.5500 - accuracy:
0.7812 - val_loss: 0.9139 - val_accuracy: 0.6750
Epoch 46/50
13/13 [=====] - 1s 52ms/step - loss: 0.5622 - accuracy:
0.7775 - val_loss: 0.8670 - val_accuracy: 0.6667
Epoch 47/50
13/13 [=====] - 1s 53ms/step - loss: 0.5550 - accuracy:
0.7700 - val_loss: 0.9434 - val_accuracy: 0.6667
Epoch 48/50
13/13 [=====] - 1s 55ms/step - loss: 0.5509 - accuracy:
0.7887 - val_loss: 0.9269 - val_accuracy: 0.7083
Epoch 49/50
13/13 [=====] - 1s 53ms/step - loss: 0.5649 - accuracy:
0.7775 - val_loss: 0.9939 - val_accuracy: 0.5750
Epoch 50/50
13/13 [=====] - 1s 52ms/step - loss: 0.5454 - accuracy:
0.7788 - val_loss: 0.8716 - val_accuracy: 0.7000

```

## 0.2 Evaluation and Testing

```
[128]: model.save('model/new_dataset/new_dataset_mood_classification_vgg16_5.h5')
```

```
[ ]: # Save the history to compare with other model
# with open("pkl/new_dataset/new_dataset_mood_classification_vgg16_5.pkl",
# ↪ "wb") as f:
#     pickle.dump(history.history, f)
```

```
[3]: model = load_model('model/new_dataset/new_dataset_mood_classification_vgg16_5.
↪ h5')
```

```
[4]: # Load the history
with open("pkl/new_dataset/new_dataset_mood_classification_vgg16_5.pkl", "rb")
↪ as f:
    history = pickle.load(f)
```

```
[5]: acc = history['accuracy']
val_acc = history['val_accuracy']
loss = history['loss']
val_loss = history['val_loss']


epochs = range(len(acc))

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
```


```
plt.legend()

plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()
```







```
[6]: def plot_confusion_matrix(cm, classes,
                                normalize=False,
                                title='Confusion matrix',
                                cmap=plt.cm.Blues):

    plt.figure(figsize = (6,6))
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=90)
    plt.yticks(tick_marks, classes)
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    thresh = cm.max() / 2.
    cm = np.round(cm,2)
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
                  horizontalalignment="center",
                  color="white" if cm[i, j] > thresh else "black")
```


```
plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.show()
```

```
[17]: predictions = model.predict(X_test)
y_pred = [np.argmax(probas) for probas in predictions]

y_test_class_indices = np.argmax(y_test, axis=1)
accuracy = accuracy_score(y_test_class_indices, y_pred)
print('Test Accuracy = %.2f' % accuracy)

confusion_mtx = confusion_matrix(y_test_class_indices, y_pred)
cm = plot_confusion_matrix(confusion_mtx, classes = list(test_labels.items()),
    ↪normalize=False)
```

```
4/4 [=====] - 40s 2s/step
Test Accuracy = 0.59
```



```
[ ]: with open("pkl/face_recognition_model.pkl", "rb") as f:
    face_recognizer = pickle.load(f)
    with open("pkl/face_embeddings.pkl", "rb") as f:
        embeddings_data = pickle.load(f)
    embeddings, labels = embeddings_data

mood_model = load_model("model/mood_classification_vgg16.h5")

face_cascade = cv2.CascadeClassifier(cv2.data.haarcascades +
    ↪ "haarcascade_frontalface_default.xml")

emotion_labels = ["angry", "happy", "sad", "confused"]
```

```

webcam_capture = cv2.VideoCapture(0)

while True:
    ret, frame = webcam_capture.read()
    if not ret:
        break
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray, 1.3, 5)

    for (x, y, w, h) in faces:
        face = frame[y: y+h, x: x+w]

        label_text = "Unknown"

        try:
            embedding = DeepFace.represent(face, model_name="Facenet",
            enforce_detection=False)[0]["embedding"]
            probs = face_recognizer.predict_proba([embedding])[0]
            pred_label = face_recognizer.predict([embedding])[0]
            face_conf = probs[1] if pred_label == 1 else probs[0]

            if pred_label == 1:
                try:
                    face_resized = cv2.resize(face, (48, 48))
                    face_array = img_to_array(face_resized) / 255.0
                    face_array = np.expand_dims(face_array, axis=0)

                    prediction = mood_model.predict(face_array, verbose=0)[0]
                    emotion = emotion_labels[np.argmax(prediction)]
                    emotion_conf = np.max(prediction)

                    label_text = f"{emotion} ({emotion_conf*100:.1f}%)"
                except:
                    label_text = "(Error)"

            else:
                label_text = f"Unknown ({face_conf*100:.1f}%)"

        except:
            label_text = "Recognition Error"

    cv2.rectangle(frame, (x, y), (x+w, y+h), (0, 255, 0), 2)
    cv2.putText(frame, label_text, (x, y-10),
                cv2.FONT_HERSHEY_SIMPLEX, 0.8, (0, 255, 0), 2)

    cv2.imshow("Face Recognition & Mood Detection", frame)

```

```
    if cv2.waitKey(1) & 0xFF == ord('q'):  
        break  
  
webcam_capture.release()  
cv2.destroyAllWindows()
```